

Fail Faster

ANONYMOUS AUTHOR(S)

[TR: Abstract needs a lot of work.] Property-based testing (PBT) is a software testing approach that verifies a system against a large suite of automatically-generated inputs. Since testing pipelines typically run under strict time constraints, PBT generators must produce inputs as quickly as possible to maximize the likelihood of finding bugs in the time available. However, existing PBT libraries often prioritize generality at the cost of performance. We introduce `waffle_house`, a high-performance generator library that uses staging, a light-weight compilation technique, to eliminate many common generator abstractions. To evaluate `waffle_house`, we design a novel benchmarking methodology that compares generators based on program equivalence, isolating performance improvements from differences in input distribution. Using this methodology, we compare `waffle_house` to a leading generator library, and, through extensive evaluation over a diverse range of generators, demonstrate that `waffle_house` significantly improves generation speed and resource efficiency while matching `base_quickcheck`'s expressiveness.

Additional Key Words and Phrases: Property-based testing, Generators, Staging, Meta-programming

1 Introduction

[JWC: This paper is about “Strict functional languages like OCaml and Scala”, but we’ll focus on OCaml for presentation.]

[BCP: The current introduction says, roughly, “We took a look at `base.quickcheck`, noticed some inefficiencies, applied known techniques from staged metaprogramming to eliminate them, and observed that speed went up.” I think we can make much stronger claims, but I’m not sure about exactly how strong or what belongs in the foreground...]

- One of the first serious uses of staging in anger?
- A use of staging that is novel / challenging in itself, in some way?
- An analysis of sources of inefficiency across a range of PBT frameworks (and a solution that applies to many of them)?
- A usable tool that addresses and overcomes some significant implementation challenges?
- Careful measurements showing where the sources of inefficiency are in existing PBT tools, and which ones matter most?
- (Your claim here...)?

]

[JWC: Agree: I think something like the 3rd is the best reframing. Something like: “We demonstrate that, across languages, monadic PBT generator DSLs can have a significant performance overhead, and present a cross-language technique for eliminating the overhead while preserving the idiomatic abstraction.”]

[JWC: People have been using staged metaprogrammign to eliminate abstraction overheads in parsing for a long time, we turn that lens on generation (Which looks like parsing)]

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

Conference acronym 'XX, Woodstock, NY

© 2025 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 978-1-4503-XXXX-X/2025/02

<https://doi.org/XXXXXX.XXXXXXXX>

50 [JWC: Another angle that I'd argue needs to come out more is that we demonstrate that engineering PBT libraries with performance in mind has significant impact on testing power.]

51 [JWC: Let's lead on with the quote from PBT in practice: performance engineering matters.
52 De-emphasize the abstraction overhead framing. Eliminating abstraction overhead is important, so
53 is using the fastest RNG you can possibly use.]

54 [JWC: "Property based testing is a race against time..."]

55 [TR: Property-based testing (PBT) uses generators of random data to produce inputs to a system
56 under test, which is then checked against a set of properties the system is expected to satisfy. In
57 practice, PBT is used as a lightweight, ad hoc approach to software correctness, aimed at detecting
58 bugs within a short testing window—typically 50 milliseconds to 30 seconds. Given this practical
59 constraint, it is imperative that inputs be generated as fast as possible; in other words, generators
60 must be treated as performance-sensitive code.]

61 [TR: The locus for improvement in this area is *generator libraries*: if programmers are provided
62 with the tools to write fast generators, then they do not themselves need to be performance experts.
63 However, existing libraries fail to achieve good performance due to two major sources of inefficiency.
64 The first is *random sampling*, or the process of selecting random values. The first is random sampling,
65 or the process of selecting random values. The second is abstraction: while generator libraries
66 are designed for flexibility and expressiveness, this comes at the cost of performance, resulting in
67 frequent boxing and unboxing and unnecessary allocations (which lead to costly garbage collection
68 pauses). In strict functional languages, abstractions like monadic bind introduce function call
69 overhead, requiring each step in a generator sequence to allocate closures and thread state through
70 nested function applications.]

71 Property-based testing (PBT) is a widely-used testing framework[HG: PBT isn't a framework, it's
72 a "technique" or "approach" – Base_QuickCheck is a framework] consisting of two key components:
73 a set of *properties* that a system must satisfy, and a way of producing a large number of *inputs* to that
74 system. In contrast to traditional unit testing, where test inputs are hand-written, PBT users create
75 a random input *generator* that automatically produces inputs. Sometimes, this process can itself
76 be automated through the use of PBT libraries that synthesize generators from type definitions;
77 however, when inputs are required to maintain invariants not communicated by the type, users
78 must write a custom generator (or filter out a broad class of invalid inputs).

79 [HG: I think this paragraph buries the lede. I'd start the paragraph talking about performance and
80 citing the 50ms number, and then jump right into what's slow.] Writing "good" generators—ones
81 that successfully uncover bugs in the system under test—is challenging. Much effort has been dedicated
82 to this problem, often by developing sophisticated domain-specific languages for specifying
83 constrained generators. However, these approaches do not always reflect the needs of real PBT
84 users. Under ideal conditions, a PBT suite would continue to run until it stopped finding bugs,
85 real-world usage patterns are very different; a study of expert PBT users reported time budgets of
86 between 50 milliseconds and 30 seconds for their test suites [3]. Consequently, generation time is a
87 significant factor in PBT efficacy—that is, if a generator can produce inputs twice as fast, it has
88 twice as many chances within a given time bound to find a bug. Therefore, it is imperative that
89 PBT libraries are built with performance in mind.

90 To better understand and address performance challenges in the landscape of PBT generators,
91 we explore popular libraries across several widely-used programming languages.[HG: Again, this
92 sentence also starts the paragraph too slowly; rather than focus on our exploration (which isn't
93 really the point of the paper) lead with the result – that abstraction overhead, boxing/unboxing,
94 and allocations slow down QuickCheck ports] We find that the flexibility of these libraries comes
95 at a cost to performance, introducing abstraction overhead, frequent boxing and unboxing, and
96 avoidable allocations (which lead to costly garbage collection pauses). We focus primarily on
97

99 OCaml's `base_quickcheck`, selected for its efficiency relative to other OCaml PBT libraries and
100 its integration with industrial build systems [HG: Let's not try to justify our choice of OCaml
101 `base_quickcheck` here — we should just talk about strict functional languages and justify this
102 choice later], where it is expected to run within tight time constraints. These inefficiencies raise
103 a fundamental question: how can PBT libraries produce code that is both efficient *and* general?
104 By analyzing `base_quickcheck` and similar libraries in other languages [TR: Rust? More?], we
105 provide a broader perspective on common performance challenges in PBT and explore potential
106 strategies for more efficient generator implementations.

107 Our solution is `waffle_house`, a generator library [HG: Can we say “an approach to PBT generator
108 libraries” or something more general like that?] that preserves `base_quickcheck`'s functionality
109 while improving performance through *multi-stage programming* (staging). Staging is a lightweight,
110 domain-specific compilation technique that allows us to completely eliminate the runtime overhead
111 of many common generator abstractions. [JWC: This relies on the fact that the code of generators is
112 known statically at compile time] For instance, in `waffle_house`, monadic operations are zero-cost.
113 `waffle_house`'s staged eliminates unnecessary allocations, generates static control structures, and
114 leverages unboxed integer libraries to achieve performance improvements over `base_quickcheck`,
115 making it well-suited to real-world PBT applications.

116 We [HG: also] present a new methodology for evaluating the relative effectiveness of PBT gener-
117 ators. [HG: Might be an over-claim, we'll need to write this carefully] Prior work has compared
118 generators based on bug-finding ability: generators are considered equally effective if they find
119 bugs in a system at the same rate over a large number of trials [5]. However, because `waffle_house`
120 preserves the semantics of `base_quickcheck`, a `waffle_house` generator produces exactly the
121 same inputs as its `base_quickcheck` counterpart. By comparing generators based on *program*
122 *equality*, we ensure that any speedups in bug-finding ability stem from performance enhancements
123 rather than variations in input distribution. This approach establishes a foundation for benchmark-
124 ing optimizations that enhance performance without compromising expressiveness. [HG: I'm not
125 totally convinced by this paragraph. Here's how I'd explain this:

126 In order to justify the switching cost associated with using `waffle_house`, we designed `waffle_house`
127 to be incredibly easy to benchmark: most PBT generator optimizations make changes up to *distribu-*
128 *tional equivalence*, allowing the precise sampling order to differ as long as the probability
129 distribution is maintained, but `waffle_house`'s optimizations are correct up to *program equivalence*.
130 This means that a `waffle_house` generator produces precisely the same values in the same order
131 as the original generator, only faster. With this setup, benchmarking becomes significantly easier
132 than for prior work—our measurements have incredibly low variance, and performance gains from
133 `waffle_house` translate directly to improved bug-finding speed.]

134 Finally, we conduct an extensive evaluation, implementing both type-derived and custom gener-
135 ators in `waffle_house` and `base_quickcheck`. Our generators produce a diverse range of inputs,
136 including recursive data structures, lambda calculus terms, and regular expressions. We benchmark
137 generation time, resource usage, and time-to-failure for each pair of staged and unstaged generators,
138 demonstrating that `waffle_house` achieves significant gains in speed and resource.

139 In summary, we dramatically improve the bug-finding effectiveness of PBT generators by opti-
140 mizing them for speed. Specifically, we make the following contributions:

141

142

143 (1) An empirical analysis of the sources of inefficiency in PBT generators. [JWC: And an

144 argument that PBT generator libraries should be engineered for performance] [HG: Let's
145 talk about this, I'm not sure I buy it]

146

147

- 148 (2) A library, `waffle_house`, which offers efficient [JWC: performant?] generator **functions → combinators** through the use of multi-stage programming; this is the first known application
 149 of staging to PBT.[HG: Let's also talk about this]
 150
 151 (3) The insight that generators should be compared by program equality
 152 (4) An evaluation demonstrating the improved performance of generators constructed using
 153 our library in a controlled comparison.

154 [JWC: Things to ensure people come away with answers to:

- 155 • Is there generality?
 156 • Is this important?
 157 • Did they actually solve the problem?

159]

160 2 Background

161 [JWC: In this paper, we'll use OCaml to illustrate the ideas, but the concepts are portable to other
 162 languages: see Section 2.3]

164 2.1 Property-Based Testing

166 [JWC: Usual spiel about pbt and generators: here's a monadic generic generator library (but we're
 167 using BQ to illustrate), it has return and bind, this is what they're for.] [JWC: You often want
 168 generators to generate for sparse preconditions so you can exercise a property reasonably well]

```
169 module Bq = struct
170   type 'a t = int -> SR.t -> 'a
171
172   let return (x : 'a) : 'a t = fun size seed -> x
173
174   let bind (g : 'a t) (f : 'a -> 'b t) : 'b t =
175     fun size seed ->
176       let x = g size seed in
177         f x) size seed
178
179   let gen_int (lo : int) (hi : int) : int t =
180     fun size seed -> SR.int seed lo hi
181 end
```

181 [JWC: `gen_int` here is really `int_uniform_inclusive`. I also think we should completely omit
 182 named arguments in the paper so we don't confuse with splices.]

183 [JWC: Here's an example of a generator: you can sample from it and it gives you pairs of ints,
 184 one less than the other.]

```
186 let int_pair : (int * int) Bq.t =
187   let%bind x = Bq.gen_int 0 100 in
188   let%bind y = Bq.gen_int 0 x in
189   return (x,y)
```

190 [JWC: This desugars to...]

```
191 let int_pair : (int * int) Bq.t =
192   Bq.bind (Bq.gen_int 0 100) (fun x ->
193     Bq.bind (Bq.gen_int 0 x) (fun y ->
194       Bq.return (x,y)
195     ))
196
```

197 [JWC:

- 198 • PBT generator libraries include more functions than just the monad interface and sampling,
libraries include helper functions.
- 200 • With `weighted_union`, to make a weighted choices, and `fixed_point` to define recursive
generators.
- 202 • Also `size` and `with_size` to adjust the size parameter — these are useful to ensure we
203 terminate with recursive generators.

204]

205 [JWC: This section needs to explain approximately how weighted union works — compute a
206 sum of the weights, sample between 0 and the sum, pick the first element where the partial sums
207 to the left exceeds the sampled number. (This has been written about a billion times, cite one of the
208 papers about this.)]

```
210 let tree_of g = fixed_point (fun rg ->
211   let%bind n = size in
212   weighted_union [
213     (1, return E);
214     (n,
215      let%bind x = g in
216      let%bind l = with_size (n / 2) rg in
217      let%bind r = with_size (n / 2) rg in
218      return (Node (l,x,r))
219    )
220  ]
221 )
```

222 [JWC: This generator generates binary trees using all the combinators, yay. Explain em.]
223 [JWC: Explain effect ordering here!!]

2.2 Multi-Stage Programming

226 [JWC: I think this section should actually go later.]

227 [JWC:

- 228 • In Multi-stage programming (also known more simply as staging), programs execute over
the course of multiple stages, with each stage producing the code to be run in the next.
- 230 • This is an old idea, with roots going back to quasiquotation in LISP, and picking up steam
again in the 1990s with MetaML [CITE].
- 232 • For the purposes of this paper, we will only consider two stages: compile time and run
time. Staged programs thus execute twice: once at compile time, which produces more code,
which is then compiled and run at run time.
- 234 • Staging has many applications, but chief among them is its use for *optimization*. We can
write programs such that during the compile time stage, they are partially evaluated to
eliminate abstraction overhead.
- 236 • Many languages have some degree of staging functionality, either provided as a library
[CITE] or build directly into the language's implementation [CITE].
- 238 • For this paper, we use MetaOCaml [CITE] for OCaml staging, but all of the staging concepts
are portable to any other language with multi-stage programming functionality.

243]

244 [JWC: This section is going to need work.]

245

246 MetaOCaml's staging functionality is exposed through a type '`a code`'. A value of type `t code`
 247 at compile time is a (potentially open) OCaml term of type `t`.

248 Values of code type are introduced by *quotes*, written `.<...>..`. Brackets delay execution of a
 249 program until run time. For example, the program `.< 5 + 1 >.` has type `int code`. Note that this
 250 is not the same as `.< 6 >..` Because brackets delay computation, the code is not *executed* until the
 251 next stage (run time). Values of type `code` can be combined together using *escape*, written `.~(e)`
 252 (or just `.~x`, when `x` is a variable). Escape lets you take a value of type `code`, and "splice" it directly
 253 into a quote. For example, this program `let x = .<1 * 5>. in .< .~(x) + .~(x) >.` evaluates
 254 to `.<(1 * 5) + (1 * 5)>..` MetaOCaml ensures correct scoping and macro hygiene, ensuring
 255 that variables are not shadowed when open terms are spliced.

256 The power of staging away abstraction overheads comes from defining functions
 257 that accept and return code values. A function `f : 'a code -> 'b code` is a function that takes
 258 a program computing a run-time '`a`' and transforms it into a program computing a run-time '`b`'. In
 259 particular, because `f` itself runs at compile time, the abstraction of using `f` is necessarily completely
 260 eliminated by run time. A code-transforming function '`'a code -> 'b code`' can also be converted
 261 *code for a function* — a value of type `('a -> 'b) code` — with the following program:

262 `let eta (f : 'a code -> 'b code) : ('a -> 'b) code = .<fun x -> .\~(f .<x>.)>.`

263 This program is known as "the trick" in the partial evaluation and multi-stage programming
 264 literature. [CITE] It returns a code for a function that takes an argument `x`, and then it splats in the
 265 result of calling `f` on just the quoted `x`. [JWC: Does this make any sense?]

266 The trick is best illustrated by an example. The following program reduces to `.< fun x -> (1`
 267 `+ x) mod 2 == 0 >.`

268 `let is_even x = .< .~x mod 2 == 0 >. in`
 269 `let succ x = .< 1 + .~x >. in`
 270 `eta (succ . is_even)`

271 By composing the two code-transforming functions together at compile time, and only then turning
 272 them into a run-time function, the functions are fused together... [JWC: words]. This is the basis of
 273 how staging is used to eliminate the abstraction of DSLs (like a generator DSL). By writing DSL
 274 combinators as compile-time functions — and only calling `eta` at the end on the completed DSL
 275 program — we can ensure that any overhead of using them is eliminated by run time.

2.3 Other Languages and Libraries

276 [JWC: TODO: Move this section later (cf conversation in WH meeting 3/7/25)]

277 [TR: Here's where we talk about what's going on in the world, and ultimately make the argument
 278 that `base_quickcheck` is the best tool to reproduce.]

279 [JWC: Note that this explanation require some prior note of exactly *why* BQ is slow... i.e. the
 280 abstraction overhead of the library is high.]

281 [JWC:

- 282 • Scala. Functional abstractions like QC generators are known to be costly in Scala, that's why
 283 they have LMS (in Scala 2, and Macros in Scala 3). Example: parser combinators ("On Staged
 284 Parser Combinators for Efficient Data Processing"), functional data structures ([Link](#)), web
 285 programming ("Efficient High-Level Abstractions for Web Programming"). All of this should
 286 still work in scala. Could easily be incorporated into ScalaCheck, with minor modification:
 287 ScalaCheck uses a state monad to thread around the seed, instead of a stateful one (like BQ),
 288 or a splittable one (like Haskell). So you have to adapt to that. But same diff.
- 289 • Python. Maybe could do it?? Hypothesis + MacroPy

- Haskell: GHC does a lot of these optimizations already, since the code is pure. Since QC generators are relatively small programs, GHC has little trouble specializing them. Of course, this is not guaranteed. A version of this idea can easily be ported to the original QC with template haskell, to guarantee the highest-performance generators.
- Rust: Not GC'd, so no alloc overhead but bind'd generators still dispatch through runtime data.

]

3 Sources of Inefficiency in Monadic Generator Libraries

[JWC: Why is a monadic generator library slow?] [JWC: NOTE: We can simplify this further by getting rid of the size parameter. Make the story even cleaner, I think!]

[JWC: IMPORTANT: We are using OCaml because we have to pick a syntax for this section, but these abstraction overheads exist in other languages – at least scala, rust.]

3.0.1 Monadic DSL Abstraction Overhead. [JWC:

- It's been long-known that clean functional abstractions have a runtime overhead (this should be familiar by this time in the paper).
- How does (simplified) BQ work?
 - The basic generator type: `'a generator = int -> SR.t -> 'a. Size and random seed to deterministic value. (note: SR.t is a mutable seed)`
 - This gets a monad instance in the obvious way (show code).
 - Also show the code for `int`, how it calls the underlying SR function.
- Show benchmarks of the running example, versus the version where you inline everything.
- Let's look at the running example: inline it all the way.

]

```

321 let int_pair : (int * int) Bq.t =
322   let%bind x = (Bq.gen_int 0 100) in
323   let%bind y = (Bq.gen_int 0 x) in
324   Bq.return (x,y)
325
326 let int_pair_inlined : (int * int) Bq.t
327   fun sr ->
328     let x = Splittable_random.int sr ~lo:0 ~hi:100 in
329     let y = Splittable_random.int sr ~lo:0 ~hi:x in
330     (x,y)

```

[JWC: Show benchmark difference between these two generators: the benchmark is in waffle-house/handwritten. It's about 2x. (see comment in latex here.)]

[JWC: The overhead of the abstraction is 2x. The reality is actually worse: actual base-quickcheck includes even more indirection in its type.]

[JWC: The native code OCaml compiler, even with `-O3`, fails to specialize this code and eliminate the overhead of this abstraction – inlining all of the function definitions and then performing beta-reductions speeds up sampling by a factor of 2.]

```

338 def intPair : Gen[(Long,Long)] = for {
339   x <- Gen.choose(0,1000)
340   y <- Gen.choose(0,x)
341 } yield (x,y)
342
343

```

```

344 def intPairInlined : (Gen.Parameters, Seed) => (Option[(Long,Long)],Seed) = {
345   (p,seed) =>
346     val (x,seed2) = chLng(0,1000)(p,seed)
347     x match {
348       case None => (None,seed2)
349       case Some(x) =>
350         val (y,seed3) = chLng(0,x)(p,seed2)
351         y match {
352           case None => (None,seed)
353           case Some(y) => (Some(x,y),seed3)
354         }
355     }
356 }
```

[JWC: Same story in scala. Approximately 2x difference for the same generator (though an order of magnitude slower overall than bq)... (scala uses a slightly different generator type, but the principles are the same) This doesn't even account for more inlining we could do to eliminate the boxing/unboxing of results. (Scalacheck generators can fail, though this adds significant overhead)]

[JWC: The problem is plain: while the monadic abstraction that PBT libraries like base quickcheck provide are indispensable for writing idiomatic generators, the performance overhead of using them is dramatic. Modern compilers use heuristics to decide when to specialize and inline code [CITE], and these heuristics cannot guarantee that generator abstractions are zero-cost. These heuristics are also necessarily conservative about inlining and specializing recursive functions, which all generators of recursive data types must be.]

[TR: Each of these consist of an explanation of the problem and pseudocode outlining the solution.]

[JWC:

- While it's the “correct” abstraction for generators, using a monadic interface prevents the compiler from specializing generator code. Using monadic bind and return obscures the control and data flow of a generator from the compiler. This is compounded when handling recursive functions.
- In cases where the compiler cannot statically eliminate them, running a monadic bind allocates a short-lived closure: we allocate a closure for the continuation, and then immediately jump into it.
- Each closure allocation is relatively cheap, but doing lots of allocation in a generation hot loop adds up fast.
- Similar things have been noticed about monadic parsing DSLs in the past.

]

3.0.2 *Combinator Abstraction Overhead*. [JWC: Aside from the usual fact that the compiler can't “see through” the abstraction boundary to specialize, combinators like union and weighted union necessarily allocate a lot in the generation hot path. Even without the dumb BQ array thing, to call union and weighted union, the types mean that you have to allocate a list. These allocations are extremely costly, and happen every sample from the generator. If the generator is recursive (like the tree generator), they occur at each recursive call.]

[JWC: In practice, the possible options are almost always statically known: in the binary tree generator from earlier, we can tell from the text of the code what the weights are, what the generators are, and how many there are. Put more simply, you basically always call weighted

392

union with an *explicit list*, which is then immediately traversed by the weighted union. For this reason, you should not have to incur the cost of allocating this list at runtime. However, almost no compilers (GHC is a notable exception) can figure this out and eliminate the list [CITE](list fusion).]

3.0.3 *Choice of RNG.* [JWC: We really need to have explained the effect ordering equivalence aspect of the evaluation before here.]

The core of any PBT generator library is a (pseudo) random number generator. Different PBT libraries use a wide range of different RNGs. Following the original Haskell QuickCheck, base_quickcheck uses the SplitMix algorithm [CITE], implemented (statefully) as an OCaml library called `Splittable_random`. Meanwhile, ScalaCheck uses the JSF algorithm [CITE].

The RNG sits at the heart of the generation hot path. Even basic generators like generating a single int or float may sample from the RNG multiple times, and generator combinators like list generate many ints. Because of this, the speed of a single RNG call matters a great deal. Unfortunately, we argue that existing PBT libraries make relatively inefficient choices on this front, leading to worse bugfinding power than what is possible.

For example, significantly faster RNG algorithms than SplitMix or JSF exist, such as the Lehmer RNG [CITE] exist. In microbenchmarks, the Lehmer generator almost 2x as fast as SplitMix [CITE]. Moreover, PBT libraries could even consider eschewing the requirement that a source of entropy pass statistical tests like BigCrush [CITE]. This is common practice in other areas of testing already: fuzzers often simply use a buffer full of arbitrary bytes as a source of entropy [CITE]. Such approaches are also faster than algorithms like SplitMix: bumping a pointer and reading from memory (which can be pipelined trivially) will always be faster on modern CPUs than any RNG with data-dependent arithmetic instructions.

In the context of this paper, however, demonstrating the benefits of any of the above faster options is tricky. Our bug-finding evaluation's power is predicated on the fact that all of the different versions of a generator we test are extensionally identical: for a given seed and size parameter, they all produce the *same* value. Using a different random number generator or different source of entropy would break this property. In Section ??, we demonstrate that faster RNG matters for bugfinding by exploiting a “natural experiment”: OCaml’s `Splittable_random` library is slow in a way that can be improved *without* changing its extensional behavior. In particular, due to implementation details related to the OCaml garbage collector, values of the OCaml type `int64` are not machine words, but rather *pointers* to machine words. This means that *all* `int64` operations (both arithmetic and bitwise) must allocate memory cells to contain their output, which has a significant performance benefit. By building a version that uses much faster “unboxed” 64-bit integer arithmetic, we can demonstrate just how much bug-finding performance can be improved just by using a more performant RNG.

4 Eliminating Abstraction Overhead of Generator DSLs by Staging

[JWC: Emphasize that this is an *COMPLETELY EQUIVALENT drop in replacement!* We didn’t just build a different library with different distributions.]

[JWC: Usual introduction to this section, corresponding to how we talked about it in the intro. “We present a library that XYZ”.]

[HG: TODO: Signposting – I got a few paragraphs in and realized I wasn’t sure what I was reading]

442 4.1 Basic Design

443 Recall that staging a DSL involves changing the combinators to run at compile time by carefully
 444 annotating their types with codes. Deciding which types `t` can instead be `t code` – in other
 445 words, figuring out which parts of the DSL can be determined statically (and can be part of the
 446 compile-time stage), and which parts are only known dynamically (and hence must be code) – is
 447 an art known as “binding-time analysis” [CITE].

448 The crux of our binding time analysis is that the particular random seed and size parameter
 449 [`JWC: if explained`] are only known at run time (the later stage), but the code of the generator
 450 itself is known at compile time.[`HG: Are we going to talk more about how we did this analysis?`]
 451 [`JWC: Well it's sort of right here: you just think about it, it's the same thing as staging a parser`
 452 `too.`]
 453 Generators – values of type `'a Bq. t` – are always constructed statically in practice, so all of
 the combinators we use to build them can run at compile time.

454 This means our library’s generator type `'a Gen. t` should have the type `int code -> SR. t`
 455 `code -> 'a code`: a compile-time function from dynamically-known size and seed to dynamically-
 456 determined result.

457 This type, along with basic monadic generator DSL functionality can be found in Figure ???. The
 458 monadic interface is given by a return and bind, as usual. Return is the constant generator, but this
 459 time it runs at compile time. Given `cx : 'a code`, the code for a `'a`, it returns the generator which
 460 always generates that value. bind `g k` sequences generators by passing the result of running the
 461 generator `g` to a continuation `k`. However, instead of getting access to the particular value generated
 462 by `g`, the continuation `k` gets access to code for the value sampled from `g`: at compile time, we only
 463 know `g` will generate *some* `'a`, but not which one ¹. Operationally, bind takes code for the size and
 464 seed, and returns code that (1) let-binds a variable `a` to spliced-in code that runs `g`, and then (2)
 465 runs the spliced-in continuation `k`. Both function applications `g size_c random_c` and `k .<a>`.
 466 `size_c random_c` run at compile time. `Gen. int` is the generator that samples an `int` from the RNG.
 467 Given any size and random seed, it returns a code block that calls `SR. int` with that random seed.
 468 Because the lower and upper bounds might not be known at compile time – they may themselves
 469 be the results of calling `Gen. int` – the arguments `lo` and `hi` are of type `int code`, and get spliced
 470 into the code block as arguments to `SR. int`. Lastly, `to_bq` turns a staged generator into code for
 471 a normal base_quickcheck generator. This function is just a 2-argument version of “The Trick”
 472 (`eta` from Section 2.2). [`HG: TODO: Be really careful with formatting of the above paragraph. It's`
 473 `very easy for some inline code to get split weirdly over a line break or just be difficult to parse in`
 474 `general, and that could throw the reader off when the content is already pretty low-level]`

475 Returning to our running example, Figure ?? shows the `int`-pair example written with the staged
 476 `Gen. t` monad, as well as the inlined code that results from calling `Gen. to_bq` (changing some
 477 identifier names for clarity). The code generated is identical to the manually inlined version from
 478 Section 3, and of course runs equally fast.

479 [`JWC: Is there anything more we need to say here?`]

481 4.2 Staging Combinators

482 In Section 3, we noted that generator combinators like `weighted_union` must allocate lists in the
 483 hot path of the generator. Even though these lists are often small – usually at most a few dozen
 484 elements in practice – each allocation takes us closer to the next garbage collection.

487 ¹Readers familiar with OCaml may notice that `return : 'a code -> 'a Gen. t` and `bind : 'a Gen. t -> ('a code ->`

488 '`b Gen. t`) -> '`b Gen. t`

489 do not have the correct types for a monad instance, preventing us from using `let%bind` notation.

We rectify this issue in Section 4.4 by importing some clever ideas from the staging literature.

```

491 module Gen = struct
492   type 'a t = int code -> Random.t code -> 'a code
493
494   let return (cx : 'a code) : 'a t = fun size_c random_c -> cx
495
496   let bind (g : 'a t) (k : 'a code -> 'b t) : 'b t =
497     fun size_c random_c ->
498       .<
499         let a = .~(g size_c random_c) in
500           .~(k .<a>. size_c random_c)
501       >.
502
503   let int (lo : int code) (hi : int code) : int t =
504     fun size_c random_c ->
505       .< SR.int .~random_c .~lo .~hi >.
506
507   let to_bq (g : 'a code Gen.t) : ('a Bq.t) code =
508     .<
509       fun size random -> .~(g .<size> . <random>.)
510     >.
511 end

```

Fig. 1. Basic Staged Generator Library
??

```

512
513
514 let int_pair_staged : (int * int) Gen.t =
515   Gen.bind (Gen.int .<0>. .<100>.) (fun cx ->
516     Gen.bind (Gen.int .<0> cx) (fun cy ->
517       Gen.return .<(.~cx,.~cy)>.
518     )
519   )
520
521 let int_pair : (int * int) Bq.t code = Gen.to_bq int_pair_staged
522 (* .< fun size random ->
523   let x = SR.int random 0 100 in
524   let y = SR.int random 0 x in
525   (x,y)
526   >.
527 *)

```

Fig. 2. Pairs of Ints, Staged
??

This is an ideal opportunity to exercise another use of staging: compile-time specialization. Since we almost always know the particular list of choices at compile time, a staged version of `weighted_union` can generate *different code* depending on the number of generators in the union. If we use `weighted_union` on a compile-time list of generators `g1`, `g2`, and `g3`, we can emit code that picks between the generators without realizing the list at run time.

Figure 3 shows the code for such a staged `weighted_union`. Crucially, it takes a *compile-time* list `weighted_gens` of generators and weights. The weights themselves might only be known at

539

```

540 module Gen =
541 ...
542   let pick (acc : int code) (weighted_gens : (int code * 'a t) list) (size : int code) (random : Sr.t code)
543     match weighted_gens with
544     | [] -> .< failwith "Error" >.
545     | (wc,g) :: gens' ->
546       .<
547         if .~acc <= .~wc then .~(g size random)
548         else
549           let acc' = .~acc - .~wc in
550             .~(pick .<acc'>. gens' size random)
551       >.
552
553   let weighted_union (weighted_gens : (int code * 'a t) list) : 'a t =
554     let sum_code = List.foldr (fun acc (w,_) -> .<.~acc + .~w>.) .<>. weighted_gens in
555     fun size random ->
556       .<
557         let sum = .~sum_code in
558         let r = SR.int .~random_c 0 sum in
559           .~(pick .<r>. weighted_gens size random)
560       >.
561
562
563

```

Fig. 3. Staged Weighted Union

run time — it is common to use the current size parameter as a weight, for instance — so they are codes. `Gen.weighted_union` begins by computing `sum_code`, an `int code` that is the sum of the weights. Note that this happens at compile time: we fold over a list known at compile time to produce another code value. We then call `SR.int` to sample a random number r between 0 and the sum. Finally, we splice in the result of calling the helper function `pick`. `pick` produces a tree of `ifs` by again traversing the list of generators at compile time. This tree of `ifs` “searches” for the generator corresponding to the sampled value r , and then runs it. [HG: I’m not sure this helps me. I think I’d prefer a less detailed explanation of the code that conveys the intuition and let the reader actually read the code if they want to know specifics. As it stands the explanation is just too dense for me to process]

Figure 4 demonstrates a use of this staged weighted union. Given a list (in this case constant) generators with weights “the current size parameter”, 2, and 1, the generated code first computes the sum of these numbers, samples between 0 and the sum, and then traverses a tree of three `ifs` to find the correct value to return.

4.3 Let-Insertion and Effect Ordering

Careful readers might note that the definition of `bind` in Section 4.1 was more complicated than one might expect. In particular, why not define `bind` in a more standard way: as `let bind' g k = fun size random -> k (g size random) size random`, without the code block that let-binds the spliced code `.~(g size random)?`[HG: Need to spell this out — the vast majority of readers won’t have that precise question in their heads. Might even be worth comparing the two implementations here side by side] Unfortunately, using `Gen.bind'` leads to incorrect code being generated. For example, consider `Gen.bind' (Gen.int .<>. .<1>.) (fun x -> Gen.return .<(. x, . x)>.)`. This generates the run time code `fun size random -> (SR.int random 0 1, int SR.random 0 1)`, which is incorrect. This is not equivalent to the behavior of writing the

```

589 let grades : char Bq.t = Gen.to_bq (
590   Gen.bind size (fun n ->
591     Gen.weighted_union [
592       (n, Gen.return .<'a'>);
593       (.<2>, Gen.return .<'b'>);
594       (.<1>, Gen.return .<'c'>);
595     ])
596   )
597 (*
598   .< fun size random ->
599     let sum = size + 2 + 1 + 0 in
600     let r = SR.int random 0 sum in
601     if r <= size then 'a'
602     else
603       let r' = r - size in
604       if r' <= 2 then 'b'
605       else
606         let r'' = r' - 2 in
607         if r'' <= 1 then 'c'
608         else
609           failwith "Error"
610     >.
611   *)
612

```

Fig. 4. Use of Staged Weighted Union

same code with base quickcheck [JWC: Which is a property we want to hold for our eval.]. Instead of generating a single integer and returning it twice, it samples two different integers.

This problem is intimately related to nondeterministic effects in the presence of CBN/CBV evaluation ordering [CITE]. In essence, the behavior of splice $\tilde{c}x$ in a staged function $f(cx : 'a \text{ code}) = \dots$ is to *copy* the entire block of code, effects and all. To ensure that the randomness effects of the first generator are executed only once, but that the value can be used in the continuation multiple times, bind let-binds the result of generation to a variable, and then passes that to the continuation.

The library [JWC: ensure this is consistent with the way we talk about the project, cf cross-language] is carefully designed to preserve exactly the effect order of base quickcheck [JWC: and the scala version to preserve the effect ordering of ScalaCheck]. [JWC: This fact should go earlier. Not really sure about where this subsection should actually land, but it needs to be said somewhere.] [HG: +1, I think this feels sort of out of place here, but I also think moving it up would make that part harder to understand too. Is there a way to tie this in to our discussion about maintaining the precise set of choices? The two issues are different, but they're related]

4.4 CodeCPS and a Monad Instance

This is all great so far,[HG: too informal] but there's a subtle issue that prevents the version of the library design discussed so far from being used as a proper drop-in replacement for an existing generator DSL: the types of $\text{return} : 'a \text{ code} \rightarrow 'a \text{ Gen.t}$ and $\text{bind} : 'a \text{ Gen.t} \rightarrow ('a \text{ code} \rightarrow 'b \text{ Gen.t}) \rightarrow 'b \text{ Gen.t}$ aren't quite right. For the type ' $a \text{ Gen.t}$ to actually be a monad, the these types cannot mention `code`. This is not just a theoretical issue, it is a significant

usability concern: the syntactic sugar for monadic programming (`let%bind` in OCaml, `foreach` in Scala, `do` in Haskell, etc) that makes it so appealing can *only* be used if the types involved are actually the proper monad function types. [HG: A pedantic reader may ask: Are we concerned with the monad laws as well, or just the type signature?]

To support real monadic programming, we'll need to adjust the type of `'a Gen.t` slightly. An initial attempt is to try type `'a t = int code -> SR.t code -> 'a`. If we strip the code off the result type, the functions `return (x : 'a) = fun _ _ -> x` and `bind (g : 'a t) (k : 'a -> 'b t) = fun size random -> k (g size random) size random` have the proper types for a monad instance. Then, any combinators of type `'a Gen.t` before simply become `'a code Gen.t` with this new version. [JWC: do we want a different name for the first cut?]

However, this definition of bind doesn't have call-by value effect semantics, as discussed in the previous section! And because the type of `g size random` is just `'a` (not necessarily `'a code`), we cannot perform the let-insertion needed to preserve the CBV effects. To solve this problem, we turn to a classic technique from the multistage programming literature: writing our staged programs in continuation-passing style [1]. [HG: Slow down! This is all very interesting and very technical! Try making each of these sentences two sentences, add some citations, and maybe spell this out with an example] [JWC: I could, but i'm just not sure this is critical.]

In Figure 5, we follow prior work [2, 4] and define the type `'a CodeCps.t = 'z. ('a -> 'z code) -> 'z code`: a polymorphic continuation transformer with the result type always in code². The monad instance for this type is the standard instance for a CPS monad with polymorphic return type. In prior work, this type is often referred to as the “code generation” monad (and sometimes, ironically, called “Gen”). This is because a value of type `('a code) CodeCps.t` is like an “action” that generates code: `CodeCps.run` passes the continuation transformer the identity continuation to produce a `'a code`. To avoid confusion with random data generators, we refer to this type as `CodeCps.t`. Most importantly, the `CodeCps` type supports a function `let_insert`, which, given `cx : 'a code`, let-binds `let x = .~cx`, and then passes `.<x>` to the continuation. [HG: I'm getting lost in the inline code again]

We can then redefine our staged generator monad type to be `'a Gen.t = int code -> SR.t code -> 'a CodeCps.t`, as shown in Figure 5. [JWC: any types that were `'a Gen.t` before are now `'a code Gen.t`]. This gives us the best of both worlds. First, we get a monad instance for `'a Gen.t` with the correct types, which lets us use the monadic syntactic sugar of our chosen language. Moreover, we also get to maintain the correct effect ordering: effectful combinators like `Gen.int` do their own let-insertion, ensuring that a program like `Gen.bind Gen.int (fun x -> ...)` generates a let-binding for the result of sampling the RNG. For example, `bind (int .<0>. .<1>) (fun cx -> return .<(. cx, . cx)>.)` now correctly generates `.< fun size random -> let x = SR.int random 0 1 in (x,x) >..` [JWC: Should we do out the whole reduction sequence here? It might be explanatory, but it also might be boring.]

This design is less obviously correct, and does require some care. Rather than bind ensuring correct evaluation order once and for all, individual combinators must be carefully written to ensure that `'a code` values that contain effects are `let_insert`'d. In Section 5.1, we discuss how we use PBT to ensure that the OCaml is written correctly.

[JWC: ... and all of this works identically in Scala, too.]

4.4.1 The Trick at Other Types.

To write more interesting generators, we also need the ability to generate code that manipulates runtime values. For instance, consider this generator

[JWC: need a better example here]

²This is an instance of the `codensity` monad [6], a fact which deserves further investigation.

```

687 module CodeCps = struct
688   type 'a t = { cps : 'z. ('a -> 'z code) -> 'z code }
689
690   let return x = {cps = fun k -> k x}
691
692   let bind (x : 'a t) (f : 'a -> 'b t) : 'b t =
693     {cps = fun k -> x.cps (fun a -> (f a).cps k)}
694
695   let run (t : ('a code) t) : 'a code = t.cps (fun x -> x)
696
697   let let_insert (cx : 'a code) : 'a code t =
698     {cps = fun k -> k .< let x = .~cx in .~(k .<x>.) >.}
699   end
700
701 module Gen = struct
702   type 'a t = int code -> SR.t code -> 'a CodeCps.t
703
704   let return (x : 'a) : 'a t = fun _ _ -> CodeCps.return x
705
706   let bind (g : 'a t) (f : 'a -> 'b t) =
707     fun size random ->
708       CodeCps.bind (g size random) (fun x ->
709         (f x) size random
710       )
711
712   let int (lo : int code) (hi : int code) : int code t =
713     fun size random -> let_insert .< SR.int .~random .~lo .~hi >.
714
715   end
716
717   let int_or_zero : int Bq.t =
718     let%bind n = size in
719     if n <= 5 then return 0 else Bq.int
720
721   let split_bool (b : bool code) : bool Gen.t =
722     fun _ _ -> _
723
724   end
725
726
727 [JWC: Describe split — this section is approximately experts-only, but you can just point at the
728 Andras paper and the things it cites.]
```

Fig. 5. CodeCPS and The Final Gen Monad

```

729
730 4.5 Recursive Generators
731 [JWC: Generating recursive datatypes requires recursive generators!]
732 Different generator DSLs handle defining recursive generators differently. Some allow recursive
733 generators to be defined as recursive functions (or in Haskell's case, recursive values), while others
734 (like base_quickcheck) expose a fixpoint combinator to construct recursive generators. Generator
735
```

Fig. 6. ??

[JWC: Describe split — this section is approximately experts-only, but you can just point at the Andras paper and the things it cites.]

Different generator DSLs handle defining recursive generators differently. Some allow recursive generators to be defined as recursive functions (or in Haskell's case, recursive values), while others (like base_quickcheck) expose a fixpoint combinator to construct recursive generators. Generator

```

736 type 'a handle
737 val recurse : 'a handle -> 'a code Gen.t
738 val fixed_point : ('a handle -> 'a code Gen.t) -> 'a code Gen.t
739

```

Fig. 7. Staged Recursive Generator Combinator API

fixpoint combinators are similar to a usual `fix : ('a -> 'a) -> 'a combinator`, except that they operate at monadic type `fixed_point : ('a Bq.t -> 'a Bq.t) -> 'a Bq.t`. Given a step function that takes a “handle” to sample from a recursive generator call, it ties the knot and builds a recursive generator.

In our case, letting programmers define recursive generators as recursive functions is out of the question. With staged programming, recursion must be handled with care: it is far too easy to accidentally recursively define an infinite code value and have the program diverge at compile time, when trying to write a code representing a recursive program. To this end, we develop a staged recursive generator combinator³, whose API is shown in Figure 7. The code for this combinator can be found in Appendix [JWC: appendix]. The recursion API consists of an opaque type `'a handle`, and a function `recurse` to perform recursive calls. Programmers can then define recursive generators by `fixed_point`, which ties the recursive knot.

4.6 Staged Type-Derived Generators

[JWC:

- In PBT in practice, we learned that in many cases, programmers don’t even write custom generators, instead relying on type-derived generators.
- These generators just produce arbitrary values of a given type, not necessarily enforcing validity conditions.
- If
- Let’s talk about how type-deriving works. In languages with typeclasses, it works by typeclass resolution. In OCaml it works by PPX system, but the principle is the same. You define rules to go from generators of a subtypes to a generator of the larger type, and then apply those rules to build up a full generator.
- – For base types, you just call the associated generator
- For product types, you sample from all the component generators, bind the values, and then tuple up the values, and return the tuple
- For variant types, you use a weighted union to choose one of the component generators with a weighted union.
- For recursive types, wrap the whole thing in a fixedpoint and then use the recursive handle as the “component generator” for all recursive instance of the type.
- This kind of generic deriving of generators works just as well for staged generators. Just replace the standard combinators in question with staged ones!
- We built this in OCaml, but you can easily use typeclasses to do it in scala too.
- Note that this is (essentially) 3-stage metaprogramming. You’re using either the PPX mechanism or typeclass resolution to generate a bunch of staged combinator calls, which then run at compile time, which then run at run time.

] [JWC: Todo: Thia]

³In reality, we actually have a more general API that allows programmers to define *parameterized* recursive combinators, of type `'r code -> 'a code Gen.t`, for any type `'r`. See Appendix [JWC: appendix] for details.

785 4.7 Fast Random Number Generators

786 [JWC:

- 787 • Using the fastest RNG possible is important.
 788 • In OCaml, the splitable random RNG is not as fast as it could be, because of the Ints issue.
 789 • In this section, we describe how to resolve the ints issue by calling out of the FFI.
 790 • All of this can be done with Jane Street bleeding-edge compiler extensions, though these
 791 are not yet compatible with MetaOCaml.
 792 • Because different libraries

793]

794 5 Evaluation

795 5.1 Implementing and Testing Generators

796 [JWC: Talk about the diff testing to ensure staging maintains semantics, and also different RNGs
 797 have identical semantics.]

800 5.2 Benchmarking speed & resource usage

801 [JWC: NOTE: we should test generator speed across both languages, but speed -> bugfinding ability
 802 in only OCaml.] [JWC: Baseline generators to test speed in both languages:

- 803 • Single int
 804 • Pair of ints, constrained
 805 • List of ints

806]

807 5.3 Impact on bug-finding ability

808 [JWC: Staging type-derived generators (which are the most commonly used ones) can turn a
 809 timeout into a found bug!]

810 6 Conclusion & Future Work

811 7 Related Work

812 [JWC: A generator is a parser of randomness, and people have implemented lots of staged parser
 813 libraries!]

814 References

- [1] Anders Bondorf. 1992. Improving binding times without explicit CPS-conversion. In *Proceedings of the 1992 ACM Conference on LISP and Functional Programming* (San Francisco, California, USA) (*LFP '92*). Association for Computing Machinery, New York, NY, USA, 1–10. <https://doi.org/10.1145/141471.141483>
- [2] Jacques Carette and Oleg Kiselyov. 2005. Multi-stage Programming with Functors and Monads: Eliminating Abstraction Overhead from Generic Code. In *Generative Programming and Component Engineering*, Robert Glück and Michael Lowry (Eds.). Springer Berlin Heidelberg, Berlin, Heidelberg, 256–274.
- [3] Harrison Goldstein, Joseph W. Cutler, Daniel Dickstein, Benjamin C. Pierce, and Andrew Head. 2024. Property-Based Testing in Practice. In *Proceedings of the IEEE/ACM 46th International Conference on Software Engineering* (Lisbon, Portugal) (*ICSE '24*). Association for Computing Machinery, New York, NY, USA, Article 187, 13 pages. <https://doi.org/10.1145/3597503.3639581>
- [4] András Kovács. 2024. Closure-Free Functional Programming in a Two-Level Type Theory. *Proc. ACM Program. Lang.* 8, ICFP, Article 259 (Aug. 2024), 34 pages. <https://doi.org/10.1145/3674648>
- [5] Jessica Shi, Alperen Keles, Harrison Goldstein, Benjamin C. Pierce, and Leonidas Lampropoulos. 2023. Etna: An Evaluation Platform for Property-Based Testing (Experience Report). *Proc. ACM Program. Lang.* 7, ICFP, Article 218 (Aug. 2023), 17 pages. <https://doi.org/10.1145/3607860>

- 834 [6] Janis Voigtländer. 2008. Asymptotic Improvement of Computations over Free Monads. In *Mathematics of Program
835 Construction*, Philippe Audebaud and Christine Paulin-Mohring (Eds.). Springer Berlin Heidelberg, Berlin, Heidelberg,
836 388–403.

837 Received –; revised –; accepted –
838

839

840

841

842

843

844

845

846

847

848

849

850

851

852

853

854

855

856

857

858

859

860

861

862

863

864

865

866

867

868

869

870

871

872

873

874

875

876

877

878

879

880

881

882