

Fail Faster

ANONYMOUS AUTHOR(S)

[TR: Abstract needs a lot of work.] Property-based testing (PBT) is a software testing approach that verifies a system against a large suite of automatically-generated inputs. Since testing pipelines typically run under strict time constraints, PBT generators must produce inputs as quickly as possible to maximize the likelihood of finding bugs in the time available. However, existing PBT libraries often prioritize generality at the cost of performance. We introduce `waffle_house`, a high-performance generator library that uses staging, a light-weight compilation technique, to eliminate many common generator abstractions. To evaluate `waffle_house`, we design a novel benchmarking methodology that compares generators based on program equivalence, isolating performance improvements from differences in input distribution. Using this methodology, we compare `waffle_house` to a leading generator library, and, through extensive evaluation over a diverse range of generators, demonstrate that `waffle_house` significantly improves generation speed and resource efficiency while matching `base_quickcheck`'s expressiveness.

Additional Key Words and Phrases: Property-based testing, Generators, Staging, Meta-programming

1 Introduction

[JWC: This paper is about “Strict functional languages like OCaml and Scala”, but we’ll focus on OCaml for presentation.]

[BCP: The current introduction says, roughly, “We took a look at `base.quickcheck`, noticed some inefficiencies, applied known techniques from staged metaprogramming to eliminate them, and observed that speed went up.” I think we can make much stronger claims, but I’m not sure about exactly how strong or what belongs in the foreground...

- One of the first serious uses of staging in anger?
- A use of staging that is novel / challenging in itself, in some way?
- An analysis of sources of inefficiency across a range of PBT frameworks (and a solution that applies to many of them)?
- A usable tool that addresses and overcomes some significant implementation challenges?
- Careful measurements showing where the sources of inefficiency are in existing PBT tools, and which ones matter most?
- (Your claim here...)?

]

[JWC: Agree: I think something like the 3rd is the best reframing. Something like: “We demonstrate that, across languages, monadic PBT generator DSLs can have a significant performance overhead, and present a cross-language technique for eliminating the overhead while preserving the idiomatic abstraction.”]

[JWC: People have been using staged metaprogramming to eliminate abstraction overheads in parsing for a long time, we turn that lens on generation (Which looks like parsing)]

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

Conference acronym 'XX, Woodstock, NY

© 2025 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 978-1-4503-XXXX-X/2025/02

<https://doi.org/XXXXXXXX.XXXXXXX>

[JWC: Another angle that I'd argue needs to come out more is that we demonstrate that engineering PBT libraries with performance in mind has significant impact on testing power.]

[JWC: Let's lead on with the quote from PBT in practice: performance engineering matters. De-emphasize the abstraction overhead framing. Eliminating abstraction overhead is important, so is using the fastest RNG you can possibly use.]

[JWC: "Property based testing is a race against time..."]

[TR: Property-based testing (PBT) uses generators of random data to produce inputs to a system under test, which is then checked against a set of properties the system is expected to satisfy. In practice, PBT is used as a lightweight, ad hoc approach to software correctness, aimed at detecting bugs within a short testing window—typically 50 milliseconds to 30 seconds. Given this practical constraint, it is imperative that inputs be generated as fast as possible; in other words, generators must be treated as performance-sensitive code.]

[TR: The locus for improvement in this area is *generator libraries*: if programmers are provided with the tools to write fast generators, then they do not themselves need to be performance experts. However, existing libraries fail to achieve good performance due to two major sources of inefficiency. The first is *random sampling*, or the process of selecting random values. (Another sentence here?) The second is abstraction: generator libraries are implemented as embedded DSLs, which are flexible and expressive, but also introduce significant overhead that is not necessarily optimized away by the compiler. These costs compound in the aggregate, raising a fundamental question: can generator libraries produce code that is both expressive *and* efficient?]

[TR: We outline a fully general, two-stage approach to designing generator libraries, which provide high-level combinator functions with dramatically reduced overhead. First, randomness library choice matters. We demonstrate that using a faster randomness library significantly impacts generator performance by directly optimizing an existing RNG, which produces the same values as its unoptimized counterpart—but faster—thereby allowing us to compare the RNG's impact on an existing generator library. Second, we describe a method of optimizing generator libraries using *staged metaprogramming*, or "staging". Staged metaprogramming allows for code generation, which can mitigate performance challenges introduced by embedded DSLs. In particular, when a generator DSL introduces overhead for dynamic control flow, staging can specialize that control flow and compile it statically. For example, when choosing between a list of options, a staged generator can avoid allocating the list of options at runtime by expanding to an if-else. This approach is not tied to any particular PBT library—instead, it gives a recipe for improving even SOTA PBT libraries.] [TR: To demonstrate the effectiveness of our technique, we apply our technique to a leading PBT library, `base_quickcheck`, built with performance engineering in mind. Our version of this library, termed `waffle_house`, is a direct drop-in replacement for the original that is completely semantically equivalent: generators written using it produce exactly the same values as those written in the original version. However, in `waffle_house`, the abstractions provided by the generator DSL are zero-cost, and the randomness library is highly-optimized. The direct equivalence between `waffle_house` and `base_quickcheck` allows us to perform extremely fine-grained comparisons between generators—generators implemented with both versions of the library will produce exactly the same sequence of values, so differences in testing performance are directly attributable to `waffle_house`'s performance speedup.] [TR: To test `waffle_house`'s effectiveness, we run a bunch of experiments to show that `waffle_house` runs faster, allocates less, etc., than `base_quickcheck` on a variety of different workloads. We demonstrate the generality of the technique by implementing `waffle_house` in Scala, and benchmarking the same generators in the Scala port of the library. Finally, as a sanity check, we demonstrate that `waffle_house` allows you to find bugs significantly faster ...]

[TR: Introduce waffle house...] [TR: Harry's "methodology" paragraph...] Property-based testing

(PBT) is a widely-used testing framework[HG: PBT isn't a framework, it's a "technique" or "approach" — Base_QuickCheck is a framework] consisting of two key components: a set of *properties* that a system must satisfy, and a way of producing a large number of *inputs* to that system. In contrast to traditional unit testing, where test inputs are hand-written, PBT users create a random input *generator* that automatically produces inputs. Sometimes, this process can itself be automated through the use of PBT libraries that synthesize generators from type definitions; however, when inputs are required to maintain invariants not communicated by the type, users must write a custom generator (or filter out a broad class of invalid inputs).

[HG: I think this paragraph buries the lede. I'd start the paragraph talking about performance and citing the 50ms number, and then jump right into what's slow.] Writing "good" generators—ones that successfully uncover bugs in the system under test—is challenging. Much effort has been dedicated to this problem, often by developing sophisticated domain-specific languages for specifying constrained generators. However, these approaches do not always reflect the needs of real PBT users. Under ideal conditions, a PBT suite would continue to run until it stopped finding bugs, real-world usage patterns are very different; a study of expert PBT users reported time budgets of between 50 milliseconds and 30 seconds for their test suites [3]. Consequently, generation time is a significant factor in PBT efficacy—that is, if a generator can produce inputs twice as fast, it has twice as many chances within a given time bound to find a bug. Therefore, it is imperative that PBT libraries are built with performance in mind.

To better understand and address performance challenges in the landscape of PBT generators, we explore popular libraries across several widely-used programming languages.[HG: Again, this sentence also starts the paragraph too slowly; rather than focus on our exploration (which isn't really the point of the paper) lead with the result — that abstraction overhead, boxing/unboxing, and allocations slow down QuickCheck ports] We find that the flexibility of these libraries comes at a cost to performance, introducing abstraction overhead, frequent boxing and unboxing, and avoidable allocations (which lead to costly garbage collection pauses). We focus primarily on OCaml's `base_quickcheck`, selected for its efficiency relative to other OCaml PBT libraries and its integration with industrial build systems[HG: Let's not try to justify our choice of OCaml `base_quickcheck` here — we should just talk about strict functional languages and justify this choice later], where it is expected to run within tight time constraints. These inefficiencies raise a fundamental question: how can PBT libraries produce code that is both efficient *and* general? By analyzing `base_quickcheck` and similar libraries in other languages [TR: Rust? More?], we provide a broader perspective on common performance challenges in PBT and explore potential strategies for more efficient generator implementations.

Our solution is `waffle_house`, a generator library[HG: Can we say "an approach to PBT generator libraries" or something more general like that?] that preserves `base_quickcheck`'s functionality while improving performance through *multi-stage programming* (staging). Staging is a lightweight, domain-specific compilation technique that allows us to completely eliminate the runtime overhead of many common generator abstractions. [JWC: This relies on the fact that the code of generators is known statically at compile time] For instance, in `waffle_house`, monadic operations are zero-cost. `waffle_house`'s staged eliminates unnecessary allocations, generates static control structures, and leverages unboxed integer libraries to achieve performance improvements over `base_quickcheck`, making it well-suited to real-world PBT applications.

We [HG: also] present a new methodology for evaluating the relative effectiveness of PBT generators.[HG: Might be an over-claim, we'll need to write this carefully] Prior work has compared generators based on bug-finding ability: generators are considered equally effective if they find bugs in a system at the same rate over a large number of trials [5]. However, because `waffle_house` preserves the semantics of `base_quickcheck`, a `waffle_house` generator produces exactly the

same inputs as its `base_quickcheck` counterpart. By comparing generators based on *program equality*, we ensure that any speedups in bug-finding ability stem from performance enhancements rather than variations in input distribution. This approach establishes a foundation for benchmarking optimizations that enhance performance without compromising expressiveness. [HG: I'm not totally convinced by this paragraph. Here's how I'd explain this:

In order to justify the switching cost associated with using `waffle_house`, we designed `waffle_house` to be incredibly easy to benchmark: most PBT generator optimizations make changes up to *distributional equivalence*, allowing the precise sampling order to differ as long as the probability distribution is maintained, but `waffle_house`'s optimizations are correct up to *program equivalence*. This means that a `waffle_house` generator produces precisely the same values in the same order as the original generator, only faster. With this setup, benchmarking becomes significantly easier than for prior work—our measurements have incredibly low variance, and performance gains from `waffle_house` translate directly to improved bug-finding speed.]

Finally, we conduct an extensive evaluation, implementing both type-derived and custom generators in `waffle_house` and `base_quickcheck`. Our generators produce a diverse range of inputs, including recursive data structures, lambda calculus terms, and regular expressions. We benchmark generation time, resource usage, and time-to-failure for each pair of staged and unstaged generators, demonstrating that `waffle_house` achieves significant gains in speed and resource.

In summary, we dramatically improve the bug-finding effectiveness of PBT generators by optimizing them for speed. Specifically, we make the following contributions:

- (1) An empirical analysis of the sources of inefficiency in PBT generators. [JWC: And an argument that PBT generator libraries should be engineered for performance] [HG: Let's talk about this, I'm not sure I buy it]
- (2) A library, `waffle_house`, which offers efficient [JWC: performant?] generator functions \rightarrow `combinators` through the use of multi-stage programming; this is the first known application of staging to PBT.[HG: Let's also talk about this]
- (3) The insight that generators should be compared by program equality
- (4) An evaluation demonstrating the improved performance of generators constructed using our library in a controlled comparison.

[JWC: Things to ensure people come away with answers to:

- Is there generality?
- Is this important?
- Did they actually solve the problem?

]

2 Background

[JWC: In this paper, we'll use OCaml to illustrate the ideas, but the concepts are portable to other languages: see Section 2.3]

2.1 Property-Based Testing

[JWC: Usual spiel about pbt and generators: here's a monadic generic generator library (but we're using BQ to illustrate), it has return and bind, this is what they're for.] [JWC: You often want generators to generate for sparse preconditions so you can exercise a property reasonably well]

```
module Bq = struct
  type 'a t = int -> SR.t -> 'a

  let return (x : 'a) : 'a t = fun size seed -> x
```

```

197
198   let bind (g : 'a t) (f : 'a -> 'b t) : 'b t =
199       fun size seed ->
200           let x = g size seed in
201           (f x) size seed
202
203   let gen_int (lo : int) (hi : int) : int t =
204       fun size seed -> SR.int seed lo hi
205 end

```

[JWC: `gen_int` here is really `int_uniform_inclusive`. I also think we should completely omit named arguments in the paper so we don't confuse with splices.]

[JWC: Here's an example of a generator: you can sample from it and it gives you pairs of ints, one less than the other.]

```

210 let int_pair : (int * int) Bq.t =
211     let%bind x = Bq.gen_int 0 100 in
212     let%bind y = Bq.gen_int 0 x in
213     return (x,y)

```

[JWC: This desugars to...]

```

215 let int_pair : (int * int) Bq.t =
216     Bq.bind (Bq.gen_int 0 100) (fun x ->
217         Bq.bind (Bq.gen_int 0 x) (fun y ->
218             Bq.return (x,y)
219         ))

```

[JWC: Explain effect ordering here!!]

- `Bq.bind` has the CBV sampling semantics
- If you `Bq.bind` `x` and then use it twice, both uses refer to the *same* value.
- Sampling is effectful: it mutates the internal state.
- Permuting two independent binds gives you a generator that is distributionally equivalent, but not equivalent when considered up to pointwise equality as *functions* from seed to value.
- Note that it's much more sensible to compare the bugfinding performance of two function-equivalent generators:

]

[JWC:

- PBT generator libraries include more functions than just the monad interface and sampling, libraries include helper functions.
- With `weighted_union`, to make a weighted choices, and `fixed_point` to define recursive generators.
- Also `size` and `with_size` to adjust the size parameter — these are useful to ensure we terminate with recursive generators.

]

[JWC: This section needs to explain approximately how weighted union works — compute a sum of the weights, sample between 0 and the sum, pick the first element where the partial sums to the left exceeds the sampled number. (This has been written about a billion times, cite one of the papers about this.)]

```

243 let tree_of g = fixed_point (fun rg ->
244     let%bind n = size in

```

```

246   weighted_union [
247     (1, return E);
248     (n,
249       let%bind x = g in
250       let%bind l = with_size (n / 2) rg in
251       let%bind r = with_size (n / 2) rg in
252       return (Node (l,x,r))
253     )
254   ]
255 )

```

[JWC: This generator generates binary trees using all the combinators, yay. Explain em.]

2.2 Multi-Stage Programming

[JWC: I think this section should actually go later.]

[JWC:

- In Multi-stage programming (also known more simply as staging), programs execute over the course of multiple stages, with each stage producing the code to be run in the next.
- This is an old idea, with roots going back to quasiquotation in LISP, and picking up steam again in the 1990s with MetaML [CITE].
- For the purposes of this paper, we will only consider two stages: compile time and run time. Staged programs thus execute twice: once at compile time, which produces more code, which is then compiled and run at run time.
- Staging has many applications, but chief among them is its use for *optimization*. We can write programs such that during the compile time stage, they are partially evaluated to eliminate abstraction overhead.
- Many languages have some degree of staging functionality, either provided as a library [CITE] or build directly into the language's implementation [CITE].
- For this paper, we use MetaOCaml [CITE] for OCaml staging, but all of the staging concepts are portable to any other language with multi-stage programming functionality.

]

[JWC: This section is going to need work.]

MetaOCaml's staging functionality is exposed through a type `'a code`. A value of type `t code` at compile time is a (potentially open) OCaml term of type `t`.

Values of code type are introduced by *quotes*, written `.<...>..`. Brackets delay execution of a program until run time. For example, the program `.< 5 + 1 >..` has type `int code`. Note that this is not the same as `.< 6 >..`. Because brackets delay computation, the code is not *executed* until the next stage (run time). Values of type `code` can be combined together using *escape*, written `.~(e)` (or just `.~x`, when `x` is a variable). Escape lets you take a value of type `code`, and “splice” it directly into a quote. For example, this program `let x = .<1 * 5>.. in .< .~(x) + .~(x) >..` evaluates to `.<(1 * 5) + (1 * 5)>..`. MetaOCaml ensures correct scoping and macro hygiene, ensuring that variables are not shadowed when open terms are spliced.

The power of staging for optimization away abstraction overheads comes from defining functions that accept and return code values. A function `f : 'a code -> 'b code` is a function that takes a program computing a run-time `'a` and transforms it into a program computing a run-time `'b`. In particular, because `f` itself runs at compile time, the abstraction of using `f` is necessarily completely eliminated by run time. A code-transforming function `'a code -> 'b code` can also be converted *code for a function* — a value of type `('a -> 'b) code` — with the following program:


```
let eta (f : 'a code -> 'b code) : ('a -> 'b) code = .<fun x -> .\~(f .<x>.)>.
```

This program is known as “the trick” in the partial evaluation and multi-stage programming literature. [CITE] It returns a code for a function that takes an argument *x*, and then it splats in the result of calling *f* on just the quoted *x*. [JWC: Does this make any sense?]

The trick is best illustrated by an example. The following program reduces to `.< fun x -> (1 + x) mod 2 == 0 >.`

```
let is_even x = .< .~x mod 2 == 0 >. in
let succ x = .< 1 + .~x >. in
eta (succ . is_even)
```

By composing the two code-transforming functions together at compile time, and only then turning them into a run-time function, the functions are fused together... [JWC: words]. This is the basis of how staging is used to eliminate the abstraction of DSLs (like a generator DSL). By writing DSL combinators as compile-time functions — and only calling *eta* at the end on the completed DSL program — we can ensure that any overhead of using them is eliminated by run time.

2.3 Other Languages and Libraries

[JWC: TODO: Move this section later (cf conversation in WH meeting 3/7/25)]

[TR: Here’s where we talk about what’s going on in the world, and ultimately make the argument that `base_quickcheck` is the best tool to reproduce.]

[JWC: Note that this explanation require some prior note of exactly *why* BQ is slow... i.e. the abstraction overhead of the library is high.]

[JWC:

- Scala. Functional abstractions like QC generators are known to be costly in Scala, that’s why they have LMS (in Scala 2, and Macros in Scala 3). Example: parser combinators (“On Staged Parser Combinators for Efficient Data Processing”), functional data structures (Link), web programming (“Efficient High-Level Abstractions for Web Programming”). All of this should still work in scala. Could easily be incorporated into ScalaCheck, with minor modification: ScalaCheck uses a state monad to thread around the seed, instead of a stateful one (like BQ), or a splittable one (like Haskell). So you have to adapt to that. But same diff.
- Python. Maybe could do it?? Hypothesis + MacroPy
- Haskell: GHC does a lot of these optimizations already, since the code is pure. Since QC generators are relatively small programs, GHC has little trouble specializing them. Of course, this is not guaranteed. A version of this idea can easily be ported to the original QC with template Haskell, to guarantee the highest-performance generators.
- Rust: Not GC’d, so no alloc overhead but bind’d generators still dispatch through runtime data.

]

3 Sources of Inefficiency in Monadic Generator Libraries

[JWC: Why is a monadic generator library slow?] [JWC: NOTE: We can simplify this further by getting rid of the size parameter. Make the story even cleaner, I think!]

[JWC: IMPORTANT: We are using OCaml because we have to pick a syntax for this section, but these abstraction overheads exist in other languages – at least scala, rust.]

3.1 Monadic DSL Abstraction Overhead

[JWC:

- It's been long-known that clean functional abstractions have a runtime overhead (this should be familiar by this time in the paper).
- How does (simplified) BQ work?
 - The basic generator type: 'a generator = int -> SR.t -> 'a. Size and random seed to deterministic value. (note: SR.t is a mutable seed)
 - This gets a monad instance in the obvious way (show code).
 - Also show the code for int, how it calls the underlying SR function.
- Show benchmarks of the running example, versus the version where you inline everything.
- Let's look at the running example: inline it all the way.

]

```

let int_pair : (int * int) Bq.t =
  let%bind x = (Bq.gen_int 0 100) in
  let%bind y = (Bq.gen_int 0 x) in
  Bq.return (x,y)

let int_pair_inlined : (int * int) Bq.t
  fun sr ->
    let x = Splittable_random.int sr ~lo:0 ~hi:100 in
    let y = Splittable_random.int sr ~lo:0 ~hi:x in
    (x,y)

```

[JWC: Show benchmark difference between these two generators: the benchmark is in waffle-house/handwr It's about 2x. (see comment in latex here.)]

[JWC: The overhead of the abstraction is 2x. The reality is actually worse: actual base-quickcheck includes even more indirection in its type.]

[JWC: The native code OCaml compiler, even with -O3, fails to specialize this code and eliminate the overhead of this abstraction — inlining all of the functon definitions and then performing beta-reductions speeds up sampling by a factor of 2.]

```

def intPair : Gen[(Long,Long)] = for {
  x <- Gen.choose(0,1000)
  y <- Gen.choose(0,x)
} yield (x,y)

def intPairInlined : (Gen.Parameters, Seed) => (Option[(Long,Long)],Seed) = {
  (p,seed) =>
    val (x,seed2) = chLng(0,1000)(p,seed)
    x match {
      case None => (None,seed2)
      case Some(x) =>
        val (y,seed3) = chLng(0,x)(p,seed2)
        y match {
          case None => (None,seed)
          case Some(y) => (Some(x,y),seed3)
        }
    }
}

```

[JWC: Same story in scala. Approximately 2x difference for the same generator (though an order of magnitude slower overall than bq)... (scala uses a slightly different generator type, but the principles are the same) This doesn't even account for more inlining we could do to eliminate the

boxing/unboxing of results. (Scalacheck generators can fail, though this adds significant overhead)

[JWC: The problem is plain: while the monadic abstraction that PBT libraries like `base quickcheck` provide are indispensable for writing idiomatic generators, the performance overhead of using them is dramatic. Modern compilers use heuristics to decide when to specialize and inline code [CITE], and these heuristics cannot guarantee that generator abstractions are zero-cost. These heuristics are also necessarily conservative about inlining and specializing recursive functions, which all generators of recursive data types must be.]

[TR: Each of these consist of an explanation of the problem and pseudocode outlining the solution.]

[JWC:

- While it's the "correct" abstraction for generators, using a monadic interface prevents the compiler from specializing generator code. Using monadic `bind` and `return` obscures the control and data flow of a generator from the compiler. This is compounded when handling recursive functions.
- In cases where the compiler cannot statically eliminate them, running a monadic `bind` allocates a short-lived closure: we allocate a closure for the continuation, and then immediately jump into it.
- Each closure allocation is relatively cheap, but doing lots of allocation in a generation hot loop adds up fast.
- Similar things have been noticed about monadic parsing DSLs in the past.

]

3.2 Combinator Abstraction Overhead

[JWC: Aside from the usual fact that the compiler can't "see through" the abstraction boundary to specialize, combinators like `union` and `weighted union` necessarily allocate a lot in the generation hot path. Even without the dumb BQ array thing, to call `union` and `weighted union`, the types mean that you have to allocate a list. These allocations are extremely costly, and happen every sample from the generator. If the generator is recursive (like the tree generator), they occur at each recursive call.]

[JWC: In practice, the possible options are almost always statically known: in the binary tree generator from earlier, we can tell from the text of the code what the weights are, what the generators are, and how many there are. Put more simply, you basically always call `weighted union` with an *explicit list*, which is then immediately traversed by the `weighted union`. For this reason, you should not have to incur the cost of allocating this list at runtime. However, almost no compilers (GHC is a notable exception) can figure this out and eliminate the list [CITE](list fusion).]

3.3 Choice of RNG

[JWC: We really need to have explained the effect ordering equivalence aspect of the evaluation before here.]

The core of any PBT generator library is a random number generator: to generate random values, we need [HG: ...] Different PBT libraries use a wide range of different RNGs. Following the original Haskell QuickCheck implementation, `base_quickcheck` uses the `SplitMix` algorithm [CITE], implemented as a (stateful) OCaml library called `Splittable_random`. Meanwhile, `ScalaCheck` uses the JSF algorithm [CITE].

The RNG sits at the heart of the hot path. Even basic generators, like ones generating a single int or float, may sample from the RNG multiple times, and generator combinators like `list` call those basic generators. [HG: To generate a single list with 3 elements, `base_quickcheck`'s `list` generator calls the RNG ...times!] Because of this, the speed of a single RNG call matters a great deal. Unfortunately, [we argue that] existing PBT libraries make relatively inefficient choices on this front, leading to worse bugfinding power than what is possible.

[HG: This paragraph feels a bit clumsy right now. What if we started with a table of the RNGs used by a bunch of popular PBT frameworks and then argued that things like Lehmer are faster?] For example, significantly faster RNG algorithms than SplitMix or JSF exist, such as the Lehmer RNG [CITE]. In microbenchmarks, the Lehmer generator almost 2x as fast as SplitMix [CITE]. Moreover, PBT libraries could even consider eschewing the requirement that a source of entropy pass statistical tests like BigCrush[CITE]. [HG: This needs more discussion] This is common practice in other areas of testing already: fuzzers often simply use a buffer full of arbitrary bytes as a source of entropy [CITE]. Such approaches are also faster than algorithms like SplitMix: bumping a pointer and reading from memory (which can be pipelined trivially) will always be faster on modern CPUs than any RNG with data-dependent arithmetic instructions. [HG: Should we just always be doing this? I think we want to argue that this is a bridge too far because you still need to generate the buffer ahead of time and you can run out of randomness?]

[HG: Alternative framing: "Of course, simply arguing that the RNG is on the hot path does not guarantee that a faster RNG leads to measurably faster generation; for that, we need an experiment. The most obvious experiment is to simply swap out the RNG for `base_quickcheck` with a totally different, faster one. But, as discussed previously, generators with different generation orders can be difficult to compare. To get around this, we exploit a "natural experiment": ..."] In the context of this paper, however, clearly demonstrating the benefits of any of the above faster options is tricky. Our bug-finding evaluation's power is predicated on the fact that all of the different versions of a generator we test are extensionally [HG: Is this how we're explaining this? I'm not sure that's quite right] identical: for a given seed and size parameter, they all produce the *same* value. Using a different random number generator or different source of entropy would break this property. In Section ??, we demonstrate that faster RNG matters for bugfinding by exploiting a "natural experiment": OCaml's `Splittable_random` library is slow in a way that can be improved *without* changing its extensional behavior. In particular, due to implementation details related to the OCaml garbage collector, values of the OCaml type `int64` are not machine words, but rather *pointers* to machine words. This means that *all* `int64` operations (both arithmetic and bitwise) must allocate memory cells to contain their output, which has a significant performance benefit. By building a version that uses much faster "unboxed" 64-bit integer arithmetic, we can demonstrate just how much bug-finding performance can be improved just by using a more performant RNG.

4 Eliminating Abstraction Overhead of Generator DSLs by Staging and Faster RNG

[JWC: Emphasize that this is an *COMPLETELY EQUIVALENT* drop in replacement! We din't just build a different library with different distributions.]

[JWC: Usual introduction to this section, corresponding to how we talked about it in the intro. "We present a library that XYZ".]

[HG: TODO: Signposting — I got a few paragraphs in and realized I wasn't sure what I was reading]

4.1 Basic Design

Recall that staging a DSL involves changing the combinators to run at compile time by carefully annotating their types with codes. Deciding which types `t` can instead be `t code` — in other

words, figuring out which parts of the DSL can be determined statically (and can be part of the compile-time stage), and which parts are only known dynamically (and hence must be code) — is an art known as “binding-time analysis” [CITE].

The crux of our binding time analysis is that the particular random seed and size parameter [JWC: if explained] are only known at run time (the later stage), but the code of the generator itself is known at compile time. [HG: Are we going to talk more about how we did this analysis?] [JWC: Well it’s sort of right here: you just think about it, it’s the same thing as staging a parser too.] Generators — values of type `'a Bq.t` — are always constructed statically in practice, so all of the combinators we use to build them can run at compile time.

This means our library’s generator type `'a Gen.t` should have the type `int code -> SR.t code -> 'a code`: a compile-time function from dynamically-known size and seed to dynamically-determined result.

This type, along with basic monadic generator DSL functionality can be found in Figure ?? . The monadic interface is given by a `return` and `bind`, as usual. `Return` is the constant generator, but this time it runs at compile time. Given `cx : 'a code`, the code for a `'a`, it returns the generator which always generates that value. `bind g k` sequences generators by passing the result of running the generator `g` to a continuation `k`. However, instead of getting access to the particular value generated by `g`, the continuation `k` gets access to code for the value sampled from `g`: at compile time, we only know `g` will generate *some* `'a`, but not which one ¹. Operationally, `bind` takes code for the size and seed, and returns code that (1) let-binds a variable `a` to spliced-in code that runs `g`, and then (2) runs the spliced-in continuation `k`. Both function applications `g size_c random_c` and `k .<a> .size_c random_c` run at compile time. `Gen.int` is the generator that samples an `int` from the RNG. Given any size and random seed, it returns a code block that calls `SR.int` with that random seed. Because the lower and upper bounds might not be known at compile time — they may themselves be the results of calling `Gen.int` — the arguments `lo` and `hi` are of type `int code`, and get spliced into the code block as arguments to `SR.int`. Lastly, `to_bq` turns a staged generator into code for a normal `base_quickcheck` generator. This function is just a 2-argument version of “The Trick” (eta from Section 2.2). [HG: TODO: Be really careful with formatting of the above paragraph. It’s very easy for some inline code to get split weirdly over a line break or just be difficult to parse in general, and that could throw the reader off when the content is already pretty low-level]

Returning to our running example, Figure ?? shows the `int-pair` example written with the staged `Gen.t` monad, as well as the inlined code that results from calling `Gen.to_bq` (changing some identifier names for clarity). The code generated is identical to the manually inlined version from Section 3, and of course runs equally fast.

[JWC: Is there anything more we need to say here?]

4.2 Staging Combinators

In Section 3, we noted that generator combinators like `weighted_union` must allocate lists in the hot path of the generator. Even though these lists are often small — usually at most a few dozen elements in practice — each allocation takes us closer to the next garbage collection.

This is an ideal opportunity to exercise another use of staging: compile-time specialization. Since we almost always know the particular list of choices at compile time, a staged version of `weighted_union` can generate *different code* depending on the number of generators in the union.

¹Readers familiar with OCaml may notice that `return : 'a code -> 'a Gen.t` and `bind : 'a Gen.t -> ('a code -> 'b Gen.t) -> 'b Gen.t` do not have the correct types for a monad instance, preventing us from using `let%bind` notation. We rectify this issue in Section 4.4 by importing some clever ideas from the staging literature.

```

540 module Gen = struct
541   type 'a t = int code -> Random.t code -> 'a code
542
543   let return (cx : 'a code) : 'a t = fun size_c random_c -> cx
544
545   let bind (g : 'a t) (k : 'a code -> 'b t) : 'b t =
546     fun size_c random_c ->
547       .<
548         let a = .~(g size_c random_c) in
549         .~(k .<a>. size_c random_c)
550       >.
551
552   let int (lo : int code) (hi : int code) : int t =
553     fun size_c random_c ->
554       .< SR.int .~random_c .~lo .~hi >.
555
556   let to_bq (g : 'a code Gen.t) : ('a Bq.t) code =
557     .<
558       fun size random -> .~(g .<size>. .<random>.)
559     >.
560 end

```

Fig. 1. Basic Staged Generator Library
??

```

563
564 let int_pair_staged : (int * int) Gen.t =
565   Gen.bind (Gen.int .<0>. .<100>.) (fun cx ->
566     Gen.bind (Gen.int .<0> cx) (fun cy ->
567       Gen.return .<(.~cx,.~cy)>.
568     )
569   )
570
571 let int_pair : (int * int) Bq.t code = Gen.to_bq int_pair_staged
572 (* .< fun size random ->
573     let x = SR.int random 0 100 in
574     let y = SR.int random 0 x in
575     (x,y)
576   >.
577 *)

```

Fig. 2. Pairs of Ints, Staged
??

If we use weighted union on a compile-time list of generators g_1 , g_2 , and g_3 , we can emit code that picks between the generators without realizing the list at run time.

Figure 3 shows the code for such a staged weighted union. Crucially, it takes a *compile-time* list `weighted_gens` of generators and weights. The weights themselves might only be known at run time — it is common to use the current size parameter as a weight, for instance — so they are codes. `Gen.weighted_union` begins by computing `sum_code`, an `int code` that is the sum of the weights. Note that this happens at compile time: we fold over a list known at compile time to

```

589 module Gen =
590 ...
591 let pick (acc : int code) (weighted_gens : (int code * 'a t) list) (size : int code) (random : Sr.t code)
592 match weighted_gens with
593 | [] -> .< failwith "Error" >.
594 | (wc,g) :: gens' ->
595   .<
596     if .~acc <= .~wc then .~(g size random)
597   else
598     let acc' = .~acc - .~wc in
599     .~(pick .<acc'>. gens' size random)
600   >.
601
602 let weighted_union (weighted_gens : (int code * 'a t) list) : 'a t =
603   let sum_code = List.foldr (fun acc (w,_) -> .<.~acc + .~w>. ) .<0>. weighted_gens in
604   fun size random ->
605     .<
606       let sum = .~sum_code in
607       let r = SR.int .~random_c 0 sum in
608       .~(pick .<r>. weighted_gens size random)
609     >.

```

Fig. 3. Staged Weighted Union

produce another code value. We then call `SR.int` to sample a random number `r` between 0 and the sum. Finally, we splice in the result of calling the helper function `pick`. `pick` produces a tree of ifs by again traversing the list of generators at compile time. This tree of ifs “searches” for the generator corresponding to the sampled value `r`, and then runs it. [HG: I’m not sure this helps me. I think I’d prefer a less detailed explanation of the code that conveys the intuition and let the reader actually read the code if they want to know specifics. As it stands the explanation is just too dense for me to process]

Figure 4 demonstrates a use of this staged weighted union. Given a list (in this case constant) generators with weights “the current size parameter”, 2, and 1, the generated code first computes the sum of these numbers, samples between 0 and the sum, and then traverses a tree of three ifs to find the correct value to return.

4.3 Let-Insertion and Effect Ordering

Careful readers might note that the definition of `bind` in Section 4.1 was more complicated than one might expect. In particular, why not define `bind` in a more standard way: as `let bind' g k = fun size random -> k (g size random) size random`, without the code block that let-binds the spliced code `.~(g size random)?` [HG: Need to spell this out — the vast majority of readers won’t have that precise question in their heads. Might even be worth comparing the two implementations here side by side] Unfortunately, using `Gen.bind'` leads to incorrect code being generated. For example, consider `Gen.bind' (Gen.int .<0>. .<1>.) (fun x -> Gen.return .<(. x, . x)>.)`. This generates the run time code `fun size random -> (SR.int random 0 1, int SR.random 0 1)`, which is incorrect. This is not equivalent to the behavior of writing the same code with base quickcheck [JWC: Which is a property we want to hold for our eval.]. Instead of generating a single integer and returning it twice, it samples two different integers.

```

638 let grades : char Bq.t = Gen.to_bq (
639   Gen.bind size (fun n ->
640     Gen.weighted_union [
641       (n, Gen.return .<'a'>.);
642       (.<2>., Gen.return .<'b'>.);
643       (.<1>., Gen.return .<'c'>.);
644     ]
645   )
646 )
647 (*
648   .< fun size random ->
649     let sum = size + 2 + 1 + 0 in
650     let r = SR.int random 0 sum in
651     if r <= size then 'a'
652     else
653       let r' = r - size in
654       if r' <= 2 then 'b'
655       else
656         let r'' = r' - 2 in
657         if r'' <= 1 then 'c'
658         else
659           failwith "Error"
660   >.
661   *)

```

Fig. 4. Use of Staged Weighted Union

This problem is intimately related to nondeterministic effects in the presence of CBN/CBV evaluation ordering [CITE]. In essence, the behavior of splice .~cx in a staged function f(cx : 'a code) = ... is to *copy* the entire block of code, effects and all. To ensure that the randomness effects of the first generator are executed only once, but that the value can be used in the continuation multiple times, bind let-binds the result of generation to a variable, and then passes that to the continuation.

The library [JWC: ensure this is consistent with the way we talk about the project, cf cross-language] is carefully designed to preserve exactly the effect order of base quickcheck [JWC: and the scala version to preserve the effect ordering of ScalaCheck]. [JWC: This fact should go earlier. Not really sure about where this subsection should actually land, but it needs to be said somewhere.] [HG: +1, I think this feels sort of out of place here, but I also think moving it up would make that part harder to understand too. Is there a way to tie this in to our discussion about maintaining the precise set of choices? The two issues are different, but they're related]

4.4 CodeCPS and a Monad Instance

This is all great so far,[HG: too informal] but there's a subtle issue that prevents the version of the library design discussed so far from being used as a proper drop-in replacement for an existing generator DSL: the types of return : 'a code -> 'a Gen.t and bind : 'a Gen.t -> ('a code -> 'b Gen.t) -> 'b Gen.t aren't quite right. For the type 'a Gen.t to actually be a monad, these types cannot mention code. This is not just a theoretical issue, it is a significant usability concern: the syntactic sugar for monadic programming (let%bind in OCaml, foreach in Scala, do in Haskell, etc) that makes it so appealing can *only* be used if the types involved are

actually the proper monad function types. [HG: A pedantic reader may ask: Are we concerned with the monad *laws* as well, or just the type signature?]

To support real monadic programming, we'll need to adjust the type of `'a Gen.t` slightly. An initial attempt is to try type `'a t = int code -> SR.t code -> 'a`. If we strip the code off the result type, the functions `return (x : 'a) = fun _ _ -> x` and `bind (g : 'a t) (k : 'a -> 'b t) = fun size random -> k (g size random) size random` have the proper types for a monad instance. Then, any combinators of type `'a Gen.t` before simply become `'a code Gen.t` with this new version. [JWC: do we want a different name for the first cut?]

However, this definition of `bind` doesn't have call-by value effect semantics, as discussed in the previous section! And because the type of `g size random` is just `'a` (not necessarily `'a code`), we cannot perform the let-insertion needed to preserve the CBV effects. To solve this problem, we turn to a classic technique from the multistage programming literature: writing our staged programs in continuation-passing style [1]. [HG: Slow down! This is all very interesting and very technical! Try making each of these sentences two sentences, add some citations, and maybe spell this out with an example] [JWC: I could, but i'm just not sure this is critical.]

In Figure 5, we follow prior work [2, 4] and define the type `'a CodeCps.t = 'z. ('a -> 'z code) -> 'z code`: a polymorphic continuation transformer with the result type always in `code`². The monad instance for this type is the standard instance for a CPS monad with polymorphic return type. In prior work, this type is often referred to as the “code generation” monad (and sometimes, ironically, called “Gen”). This is because a value of type `('a code) CodeCps.t` is like an “action” that generates code: `CodeCps.run` passes the continuation transformer the identity continuation to produce a `'a code`. To avoid confusion with random data generators, we refer to this type as `CodeCps.t`. Most importantly, the `CodeCps` type supports a function `let_insert t`, which, given `cx : 'a code`, let-binds `let x = .~cx`, and then passes `.<x>.` to the continuation. [HG: I'm getting lost in the inline code again]

We can then redefine our staged generator monad type to be `'a Gen.t = int code -> SR.t code -> 'a CodeCps.t`, as shown in Figure 5. [JWC: any types that were `'a Gen.t` before are now `'a code Gen.t`]. This gives us the best of both worlds. First, we get a monad instance for `'a Gen.t` with the correct types, which lets us use the monadic syntactic sugar of our chosen language. Moreover, we also get to maintain the correct effect ordering: effectful combinators like `Gen.int` do their *own* let-insertion, ensuring that a program like `Gen.bind Gen.int (fun x -> ...)` generates a let-binding for the result of sampling the RNG. For example, `bind (int .<0>. .<1.>) (fun cx -> return .<(. cx, . cx)>.)` now correctly generates `.< fun size random -> let x = SR.int random 0 1 in (x,x) >..` [JWC: Should we do out the whole reduction sequence here? It might be explanatory, but it also might be boring.]

This design is less obviously correct, and does require some care. Rather than `bind` ensuring correct evaluation order once and for all, individual combinators must be carefully written to ensure that `'a code` values that contain effects are `let_insert t`'d. In Section 5.1, we discuss how we use PBT to ensure that the OCaml is written correctly.

[JWC: ... and all of this works identically in Scala, too.]

4.4.1 The Trick at Other Types. To write more interesting generators, we also need the ability to generate code that manipulates runtime values. For instance, consider this generator

[JWC: need a better example here]

[JWC: Describe split — this section is approximately experts-only, but you can just point at the Andras paper and the things it cites.]

²This is an instance of the *codensity* monad [6], a fact which deserves further investigation.

```

736 module CodeCps = struct
737   type 'a t = { cps : 'z. ('a -> 'z code) -> 'z code }
738
739   let return x = {cps = fun k -> k x}
740
741   let bind (x : 'a t) (f : 'a -> 'b t) : 'b t =
742     {cps = fun k -> x.cps (fun a -> (f a).cps k)}
743
744   let run (t : ('a code) t) : 'a code = t.cps (fun x -> x)
745
746   let let_insert (cx : 'a code) : 'a code t =
747     {cps = fun k -> k .< let x = .~cx in .~(k .<x>.) >.>}
748 end
749
750 module Gen = struct
751   type 'a t = int code -> SR.t code -> 'a CodeCps.t
752
753   let return (x : 'a) : 'a t = fun _ _ -> Codecps.return x
754
755   let bind (g : 'a t) (f : 'a -> 'b t) =
756     fun size random ->
757       CodeCps.bind (g size random) (fun x ->
758         (f x) size random
759       )
760
761   let int (lo : int code) (hi : int code) : int code t =
762     fun size random -> let_insert .< SR.int .~random .~lo .~hi >.>
763 end

```

Fig. 5. CodeCPS and The Final Gen Monad

```

766 let int_or_zero : int Bq.t =
767   let%bind n = size in
768   if n <= 5 then return 0 else Bq.int
769
770 let split_bool (b : bool code) : bool Gen.t =
771   fun _ _ -> _
772
773

```

Fig. 6. ??

4.5 Recursive Generators

[JWC: Generating recursive datatypes requires recursive generators!]

Different generator DSLs handle defining recursive generators differently. Some allow recursive generators to be defined as recursive functions (or in Haskell's case, recursive values), while others (like `base_quickcheck`) expose a fixpoint combinator to construct recursive generators. Generator fixpoint combinators are similar to a usual `fix : ('a -> 'a) -> 'a` combinator, except that they operate at monadic type `fixed_point : ('a Bq.t -> 'a Bq.t) -> 'a Bq.t`. Given a step

```

type 'a handle
val recurse : 'a handle -> 'a code Gen.t
val fixed_point : ('a handle -> 'a code Gen.t) -> 'a code Gen.t

```

Fig. 7. Staged Recursive Generator Combinator API

function that takes a “handle” to sample from a recursive generator call, it ties the knot and builds a recursive generator.

In our case, letting programmers define recursive generators as recursive functions is out of the question. With staged programming, recursion must be handled with care: it is far too easy to accidentally recursively define an infinite code value and have the program diverge at compile time, when trying to write a code representing a recursive program. To this end, we develop a staged recursive generator combinator³, whose API is shown in Figure 7. The code for this combinator can be found in Appendix [JWC: appendix]. The recursion API consists of an opaque type `'a handle`, and a function `recurse` to perform recursive calls. Programmers can then define recursive generators by `fixed_point`, which ties the recursive knot.

4.6 Staged Type-Derived Generators

[JWC:

- In PBT in practice, we learned that in many cases, programmers don’t even write custom generators, instead relying on type-derived generators.
- These generators just produce arbitrary values of a given type, not necessarily enforcing validity conditions.
- If
- Let’s talk about how type-deriving works. In languages with typeclasses, it works by typeclass resolution. In OCaml it works by PPX system, but the principle is the same. You define rules to go from generators of a subtypes to a generator of the larger type, and then apply those rules to build up a full generator.
 - For base types, you just call the associated generator
 - For product types, you sample from all the component generators, bind the values, and then tuple up the values, and return the tuple
 - For variant types, you use a weighted union to choose one of the component generators with a weighted union.
 - For recursive types, wrap the whole thing in a fixedpoint and then use the recursive handle as the “component generator” for all recursive instance of the type.
- This kind of generic deriving of generators works just as well for staged generators. Just replace the standard combinators in question with staged ones!
- We built this in OCaml, but you can easily use typeclasses to do it in scala too.
- Note that this is (essentially) 3-stage metaprograming. You’re using either the PPX mechanism or typeclass resolution to generate a bunch of staged combinator calls, which then run at compile time, which then run at run time.

] [JWC: Todo: Thia]

³In reality, we actually have a more general API that allows programmers to define *parameterized* recursive combinators, of type `'r code -> 'a code Gen.t`, for any type `'r`. See Appendix [JWC: appendix] for details.

4.7 Faster SplitMix with Unboxed int64s

[JWC: Unclear what we should call this section] As we discussed in Section 3, inefficient RNG choice is another bottleneck for finding bugs fast. While generator libraries by and large use sensible RNG algorithms, they are not explicitly chosen with performance in mind. Indeed, much faster RNGs [CITE] and fast sources of entropy [CITE] exist. [JWC: ... more here?] To demonstrate that a faster RNG can significantly impact bugfinding power, we use the natural experiment provided by OCaml's inefficient implementation of SplitMix. By replacing this slow RNG with a faster but extensionally equivalent implementation, we are able to precisely quantify the bugfinding speedup that using a RNG gives across a range of PBT scenarios.

We emphasize that while we believe that the insight that faster RNG translates to faster bug finding is a cross-language one, the specific technical contents of this section are OCaml-specific. In most other PBT frameworks [CITE], the RNG operates on by machine integers, so this *particular* inefficiency does not exist.

The precise details of how the SplitMix algorithm works [CITE] are unimportant for the present paper, but the critical component is that all of its operations are defined in terms of arithmetic bitwise operations on 64-bit integers. In OCaml, because of details related to the garbage collector, the 64-bit integer type `int64` is represented at run time as a *pointer* to an unscanned block of memory containing (among other things) a 64-bit integer [CITE]. This means that all operations that return an `int64` must allocate this block of memory. This has a significant impact on the performance of generators. A single call to one of the `base_quickcheck` library functions — like generating an integer uniformly in a range — may call the `splittable_random` algorithm multiple times. Each sample from `splittable_random` allocates 9 times [JWC: this is a call to `next_int64`], and each allocation brings us closer to the next garbage collection pause. While small allocations like these are very fast to perform and subsequently collect in OCaml,⁴ we will see in Section ?? that this can have a large performance impact on some generators that spend most of their time sampling data. To circumvent this allocation and provide an equivalent version of `splittable_random`, we reimplement SplitMix in C, and call out to it with the OCaml FFI. The C version of the library uses proper `int64_t` arithmetic, only boxing and unboxing integers at the call boundaries between OCaml and C code. Ideally in the future, one would not need to call out to C for this: the Jane Street bleeding-edge OCaml compiler has support for unboxed types [CITE], which (among other things) would let us implement a version of SplitMix that does not allocate, directly in OCaml. Unfortunately, the Jane Street branch of the compiler is incompatible with MetaOCaml, which we use to implement the metaprogramming discussed in the previous sections.

5 Evaluation

5.1 Implementing and Testing Generators

[JWC: Talk about the diff testing to ensure staging maintains semantics, and also different RNGs have identical semantics.]

5.2 Benchmarking speed & resource usage

[JWC: NOTE: we should test generator speed across both languages, but speed -> bugfinding ability in only OCaml.] [JWC: Baseline generators to test speed in both languages:

- Single int
- Pair of ints, constrained
- List of ints without bind (use map): lots of sampling, minimal binds.

⁴The OCaml GC is a generational collector [CITE], and since these allocations are small and mostly very short lived, they will all be minor allocations, never to be promoted.

- Unableled Trees of a fixed size (no weighted union) minimal sampling, lots of binds.

5.3 Impact on bug-finding ability

[JWC: Staging type-derived generators (which are the most commonly used ones) can turn a timeout into a found bug!]

6 Conclusion & Future Work

7 Related Work

[JWC: A generator is a parser of randomness, and people have implemented lots of staged parser libraries!]

References

- [1] Anders Bondorf. 1992. Improving binding times without explicit CPS-conversion. In *Proceedings of the 1992 ACM Conference on LISP and Functional Programming* (San Francisco, California, USA) (*LFP '92*). Association for Computing Machinery, New York, NY, USA, 1–10. <https://doi.org/10.1145/141471.141483>
- [2] Jacques Carette and Oleg Kiselyov. 2005. Multi-stage Programming with Functors and Monads: Eliminating Abstraction Overhead from Generic Code. In *Generative Programming and Component Engineering*, Robert Glück and Michael Lowry (Eds.). Springer Berlin Heidelberg, Berlin, Heidelberg, 256–274.
- [3] Harrison Goldstein, Joseph W. Cutler, Daniel Dickstein, Benjamin C. Pierce, and Andrew Head. 2024. Property-Based Testing in Practice. In *Proceedings of the IEEE/ACM 46th International Conference on Software Engineering* (Lisbon, Portugal) (*ICSE '24*). Association for Computing Machinery, New York, NY, USA, Article 187, 13 pages. <https://doi.org/10.1145/3597503.3639581>
- [4] András Kovács. 2024. Closure-Free Functional Programming in a Two-Level Type Theory. *Proc. ACM Program. Lang.* 8, ICFP, Article 259 (Aug. 2024), 34 pages. <https://doi.org/10.1145/3674648>
- [5] Jessica Shi, Alperen Keles, Harrison Goldstein, Benjamin C. Pierce, and Leonidas Lampropoulos. 2023. Etna: An Evaluation Platform for Property-Based Testing (Experience Report). *Proc. ACM Program. Lang.* 7, ICFP, Article 218 (Aug. 2023), 17 pages. <https://doi.org/10.1145/3607860>
- [6] Janis Voigtländer. 2008. Asymptotic Improvement of Computations over Free Monads. In *Mathematics of Program Construction*, Philippe Audebaud and Christine Paulin-Mohring (Eds.). Springer Berlin Heidelberg, Berlin, Heidelberg, 388–403.

Received –; revised –; accepted –